

The ELF Format

This assignment will make you more familiar with the organization of ELF files. You can do this assignment on any operating system that supports the Unix API (Linux Openlab machines, your laptop that runs Linux or Linux VM, and even MacOS). **You don't need to set up Xv6 for this assignment.**

For MacOS users, the support of 32 bit applications is deprecated in the latest version of your system. So if you already updated your system to MacOS Catalina or have updated your XCode then we recommend you to do the homework at the Openlab machines.

Part 1: Take a Look at ELF files

Top

Download the `main.c` , and `elf.c` and look them over. At a high level this homework asks you to implement a simple ELF loader (you will extend the `main.c` file) and use it to load a simple ELF object file (the one compiled from `elf.c`). However, before starting this task, lets make ourselves familiar with ELF files.

We provide a simple `Makefile` that compiles `elf.o` and `main` as ELF executables. Look over the `Makefile` and then compile both files by running:

```
$ make
```

Now, let's take a look at the ELF files that we just compiled. We will use the `readelf` tool:

```
$ readelf -a elf
```

ELF is the file format used for `[.o]` object files, binaries, shared libraries and core dumps in Linux. It's actually pretty simple and well thought-out.

ELF has the same layout for all architectures, however endianness and word size can differ; relocation types, symbol types and the like may have platform-specific values, and of course the contained code is architecture specific.

ELF files are used by two tools: the **linker** and the **loader**. A linker combines multiple ELF files into an executable or a library and a loader loads the executable ELF file in the memory of the process. On real operating systems, loading may require relocation (e.g., if the file is dynamically linked it has to be linked again with all the shared libraries it depends on). In this homework we will not do any relocation (it's too complicated), we'll simply load an ELF file in memory and run it.

Linker and loader need two different views of the ELF file, i.e., **they access it differently**.

On the one hand, the linker needs to know where the DATA, TEXT, BSS, and other sections are to merge them with sections from other libraries. If relocation is required the linker needs to know where the symbol tables and relocation information is.

On the other hand, the loader does not need any of these details. It simply needs to know which parts of the ELF file are code (executable), which are data and read-only data, and where to put the BSS in the memory of a process.

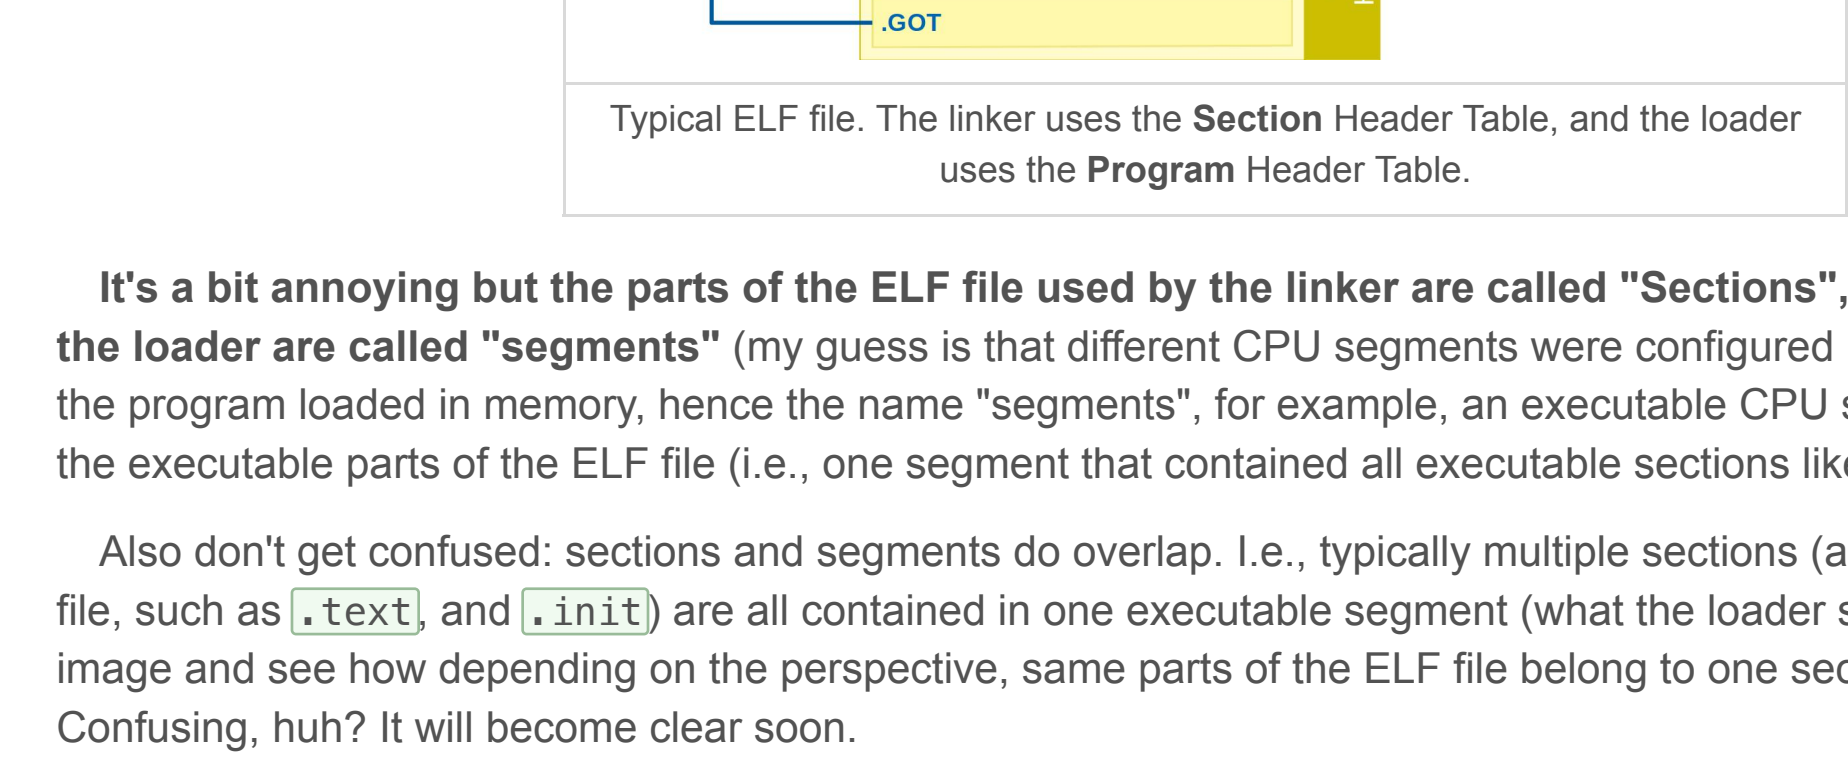
Hence, the ELF file provides two separate views on the data inside the ELF file: a **linker view** with several details, and a **loader view**, a higher level view with less details.

To provide these views each ELF file contains two tables (or arrays):

- **Section Header Table**: With pointers to **sections** to be used by the **linker**.
- **Program Header Table**: With pointers to **segments** to be used by the **loader**.

Both tables are simply arrays of entries that contain information about each part of the ELF file (e.g., where the sections/segments used by the linker/loader are inside the ELF file).

Here is a simple figure of a typical ELF file that starts with the ELF header. The header contains pointers to the locations of Section Header Table and Program Header Table within the ELF file. Then each tables have entries that point to the starting locations of individual sections and segments.



It's a bit annoying but the parts of the ELF file used by the linker are called **"Sections"**, and the parts used by the loader are called **"segments"** (my guess is that different CPU segments were configured in the past for each part of the program loaded in memory, hence the name "segments", for example, an executable CPU segment was created for the executable parts of the ELF file (i.e., one segment that contained all executable sections like `.text` and `.init`, etc.).

Also don't get confused: sections and segments do overlap. I.e., typically multiple sections (as the linker sees the ELF file, such as `.text` and `.init`) are all contained in one executable segment (what the loader sees). Check the previous image and see how depending on the perspective, same parts of the ELF file belong to one section **and** one segment. Confusing, huh? It will become clear soon.

Linking View: Section Header Table (SHT)

Top

The Section Header Table is an array in which every entry contains a pointer to one of the sections of the ELF file. Lets take a look at what inside the ELF file. Run this command, and scroll down to the **Section headers** you will see all "sections" of the ELF file that the linker can use:

```
$ readelf --section-headers elf
There are 12 section headers, starting at offset 0x458:

Section Headers:
 [Nr] Name                Type             Addr             Off             Size             ES Flg Lk Inf Al
 [ 0] .text                  PROGBITS         00000000          000000          000000          00  0  0  0
 [ 1] .note.ABI-tag          NOTE             00000168          000168          000020          00  0  0  4
 [ 2] .eh_frame             PROGBITS         00000010          00000A          000038          00  0  0  0
 [ 3] .comment              PROGBITS         00000000          00000c          00002b          01  MS  0  1
 [ 4] .debug_aranges       PROGBITS         00000000          0000e7          000020          00  0  0  1
 [ 5] .debug_info          PROGBITS         00000000          000107          000066          00  0  0  1
 [ 6] .debug_abbrev        PROGBITS         00000000          00016d          000055          00  0  0  1
 [ 7] .debug_line          PROGBITS         00000000          0001c2          000035          00  0  0  1
 [ 8] .debug_str            PROGBITS         00000000          0001f7          000106          01  MS  0  1
 [ 9] .symtab               SYMTAB           00000000          000300          0000e0          10  10  4  4
 [10] .strtab               STRTAB           00000000          0003e0          000024          00  0  0  1
 [11] .shstrtab             STRTAB           00000000          000404          000074          00  0  0  1

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
P (processor specific)
```

Since `elf.c` is a very simple program, it has only `.text` section (i.e., code of the program), a bunch of sections that contain debugging information, and a `.symtab` section that contains imported and exported symbols.

Note, there is no `.data` or `.bss` sections for global variables (there are no globals in `elf.c`).

You can experiment by adding a not initialized global variable to `elf.c`, recompile, and see the difference in the SHT. Then add an initialized global variable, recompile and check again the SHT.
Make sure of removing the global variables and recompile before continuing with the assignment.)

Moreover, since we linked `elf.c` to be a static executable that is linked to run if loaded at address `0x0` (the `Text 0` in the `Makefile`) tells the linker to relocate the executable at linking time to work at `0x0`.

```
$ cat Makefile
...
elf: elf.o
ld -m elf_i386 -N -e main -Ttext 0 -o elf elf.o
elf.o: elf.c
$(CC) -c -fno-pic -static -fno-builtin -ggdb -m32 -fno-omit-frame-pointer elf.c
```

The **symbol table** contains these symbols

```
$ readelf --syms elf
Symbol table '.symtab' contains 14 entries:
Num: Value             Size Type             Bind Vis            Ndx Name
 0: 00000000             0 NOTYPE          LOCAL DEFAULT    1
 1: 00000000             0 SECTION        LOCAL DEFAULT    2
 2: 00000010             0 SECTION        LOCAL DEFAULT    3
 3: 00000000             0 SECTION        LOCAL DEFAULT    4
 4: 00000000             0 SECTION        LOCAL DEFAULT    5
 5: 00000000             0 SECTION        LOCAL DEFAULT    6
 6: 00000000             0 SECTION        LOCAL DEFAULT    7
 7: 00000000             0 SECTION        LOCAL DEFAULT    8
 8: 00000000             0 FILE           LOCAL DEFAULT    ABS elf.c
 9: 00000048             0 NOTYPE          GLOBAL DEFAULT   2 _bss_start
11: 00000000            13 FUNC            GLOBAL DEFAULT   1 main
12: 00000048             0 NOTYPE          GLOBAL DEFAULT   2 _edata
13: 00000048             0 NOTYPE          GLOBAL DEFAULT   2 _end
```

The `main` is our function (it's `FUNC`, and `GLOBAL`), the `_bss_start`, `_edata`, and `_end` are added by the linker to mark the start and end of the BSS, TEXT, and DATA sections.

If we take a look at the `main` executable, the ELF file is more complicated.

```
$ readelf --section-headers main
There are 35 section headers, starting at offset 0x4794:

Section Headers:
 [Nr] Name                Type             Addr             Off             Size             ES Flg Lk Inf Al
 [ 0] .interp              PROGBITS         00000154          000154          000013          00  A  0  0  1
 [ 1] .note.gnu.build-id   NOTE             00000168          000168          000024          00  0  0  4
 [ 2] .gnu.hash           GNU_HASH         000001ac          0001ac          000020          04  A  5  0  4
 [ 3] .dynsym             DYNSYM          000001cc          0001cc          000090          10  A  6  1  4
 [ 4] .dynstr             STRTAB          0000025c          00025c          000008          00  0  0  1
 [ 5] .gnu.version        VERNEED         00000314          000314          000012          02  A  5  0  2
 [ 6] .gnu.version_r      VERNEED         00000328          000328          000040          00  A  6  1  4
 [ 7] .rel.dyn            REL             00000368          000368          000058          00  A  5  0  4
 [ 8] .rel.plt            REL             000003c0          0003c0          000118          00  AI  5  22  4
 [11] .init               PROGBITS         000003d8          0003d8          000023          00  AX  0  0  4
 [12] .plt               PROGBITS         00000400          000400          000040          04  AX  0  0  16
 [13] .plt.got           PROGBITS         00000440          000440          000010          00  AX  0  0  8
 [14] .text             PROGBITS         00000450          000450          000222          00  AX  0  0  16
 [15] .fini             PROGBITS         00000574          000574          000014          00  AX  0  0  4
 [16] .rodata           PROGBITS         00000688          000688          000010          00  A  0  0  4
 [17] .eh_frame_hdr     PROGBITS         00000698          000698          000034          00  A  0  0  4
 [18] .eh_frame         PROGBITS         000006cc          0006cc          0000e0          00  A  0  0  4
 [19] .init_array       INIT_ARRAY       00001ecc          000ecc          000004          04  WA  0  0  4
 [20] .fini_array       FINI_ARRAY       00001ed0          000ed0          000004          04  WA  0  0  4
 [21] .dynamic          DYNAMIC          00001ed4          000ed4          000100          00  WA  6  0  4
 [22] .got              PROGBITS         00001fd4          000fd4          00002c          04  WA  0  0  4
 [23] .data             PROGBITS         00002000          001000          000008          00  WA  0  0  4
 [24] .bss              NOBITS           00002008          001008          000004          00  WA  0  0  1
 [25] .comment          PROGBITS         00000000          001008          00002b          01  MS  0  0  1
 [26] .debug_aranges    PROGBITS         00000000          001033          000020          00  0  0  1
 [27] .debug_info       PROGBITS         00000000          001053          000554          00  0  0  1
 [28] .debug_abbrev     PROGBITS         00000000          0015a7          000132          00  0  0  1
 [29] .debug_line       PROGBITS         00000000          0016d9          00000e          00  0  0  1
 [30] .debug_str        PROGBITS         00000000          0017c7          0003b4          01  MS  0  0  1
 [31] .symtab           SYMTAB           00000000          001b7c          000480          10  32  48  4
 [32] .strtab           STRTAB           00000000          001ffc          00024f          00  0  0  1
 [33] .shstrtab        STRTAB           00000000          00224b          00013c          00  0  0  1

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
P (processor specific)
```

It contains all the section we've mentioned in class: `.text` (main code of the program), `.data` (data section for global variables), `.rodata` (data section for global read-only variables), `.bss` (uninitialized global variables), `.init` (init section to call the constructors that run before `main()`), `.got` (Global Offset Table), `.plt` (Procedure Linking Table for lazy linking of imported functions), and even the `.interp` (the section for the interpreter, i.e., the linker that links dynamically linked program before it runs, typically it's `/lib/ld-linux.so.2` on Linux systems).

Execution view: Program Header Table (PHT)

Top

The Program Header Table contains information for the kernel on how to start the program. The `LOAD` directives determine what parts of the ELF file get mapped into program memory.

Again, in our `elf` example the program header defines only two segments. And only one of them should be loaded by the operating system in memory to run.

```
$ readelf --program-headers elf
Elf file type is EXEC (Executable file)
Entry point 0x0
There are 2 program headers, starting at offset 52

Program Headers:
Type           Offset             VirtAddr          PhysAddr         FileSiz MemSiz  Flg Algn
LOAD           0x00000074         0x00000000         0x00000000        0x00048 0x00048  RW  0x4
GNU_STACK      0x00000000         0x00000000         0x00000000        0x00000 0x00000  RW  0x10

Section to Segment mapping:
Segment Sections...
00          .text .eh_frame
01
```

The only loadable section is linked to run at address `0x00000000`. We can inspect the `elf` binary with the `objdump` tool to see what is there:

```
$ objdump -d elf
elf:          file format elf32-i386

Disassembly of section .text:

00000000 <main>:
0: 55                      push    %ebp
1: 89 e5                   mov     %esp,%ebp
3: 8b 55 08                mov     0x8(%ebp),%edx
6: 8b 45 0c                mov     0xc(%ebp),%eax
9: 01 d0                   add     %edx,%eax
b: 5d                      pop     %ebp
c: c3                      ret
```

Well, no surprises: it's the main function compiled into machine code.

Putting it All Together

Top

Neither the SHT nor the PHT have fixed positions, they can be located anywhere in an ELF file. To find them the ELF header is used, which is located at the very start of the file.

The first bytes contain the elf magic `"\x7fELF"`, followed by the class ID (32 or 64 bit ELF file), the data format ID (little endian/big endian), the machine type, etc.

At the end of the ELF header are then pointers to the SHT and PHT. Specifically, the Segment Header Table which is used by the linker starts at byte 1120 in the ELF file, and the Program Header Table starts at byte 52 (right after the ELF header)

```
$ readelf --file-header elf
ELF Header:
Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
Class:                                ELF32
Data:                                  2's complement, little endian
Version:                              1 (current)
OS/ABI:                               UNIX - System V
ABI Version:                           0
Type:                                  EXEC (Executable file)
Machine:                               Intel 80386
Version:                               0x1
Entry point address:                   0x0
Start of program headers:              52 (bytes into file)
Start of section headers:              1120 (bytes into file)
Flags:                                 0x0
Size of this header:                   52 (bytes)
Size of program headers:               32 (bytes)
Number of program headers:             2
Size of section headers:               40 (bytes)
Number of section headers:             12
Section header string table index:     11
```

Finally, the entry point of this file is at address `0x0`. This is exactly what we told the linker to do — link the program to run at address `0x00000000`. And this is where the `main` function of the `elf.c` file is shown in the `objdump`.

Program Loading in the Kernel

Top

The execution of a program starts inside the kernel, in the `exec("/bin/wc",...)` system call takes a path to the executable file. The kernel reads the ELF header and the program header table (PHT), followed by lots of sanity checks.

For statically linked executables:

1. The kernel reads the PHT, and loads the parts specified in the `LOAD` directives into memory.
2. The parts specified in `LOAD` directives are loaded, control can be transferred to the entry point of the program, which is `main()`.

For dynamically linked executables:

1. The kernel reads the PHT, and loads the parts specified in the `INTERP` and `LOAD` directives into memory. Dynamically linked programs always need `/lib/ld-linux.so.2` as interpreter because it includes some startup code, loads shared libraries needed by the binary, and performs relocations.
2. Once the parts specified in `INTERP` and `LOAD` directives are loaded, control can be transferred to the interpreter.

The dynamic linker (contained within the interpreter):

1. Looks at the `.dynamic` section, whose address is stored in the PHT, and finds:
2. The `NEEDED` entries determining which libraries have to be loaded before the program can be run.
3. The `REL**` entries containing the address of the relocation tables.
4. The `VER**` entries which contain symbol versioning information.
5. Etc...

The dynamic linker loads the needed libraries and performs relocations (either directly at program startup or later, as soon as the relocated symbol is needed, depending on the relocation type).

Finally, control is transferred to the address given by the symbol `__start` in the binary. Normally some gcc/glibc startup code lives there, which in the end calls `main()`.

What do you need to do in this homework: load an ELF file

Top

While ELF might look a bit intimidating, in practice the loading algorithm is straightforward:

1. Read the ELF header.
2. One of the ELF header fields tells you the offset of the program header table inside the file.
3. Read each entry of the program header table (i.e., read each program header)
4. Each program header has an offset and size of a specific segment inside the ELF file (e.g., a executable code). You have to read it from the file and load it in memory.
5. When done with all segments, jump to the entry point of the program. (Note since we don't control layout of the address space at the moment, we load the sections at some random place in memory (the place that is allocated for us by the `mmap()` function). Obviously the address of the entry point should be an offset within that random area. Such loading will not work for a real ELF file, but ours is simple: it's statically linked, and contains the code that can run at any location in memory. So even though it's linked to run at `0x0` it will run anywhere where it was loaded.

Looks manageable. We make a couple of simplifications. First we create a very simple ELF file out of `elf.c`: it contains only one function, it has no data, and the code can be placed anywhere in memory and will run ok (it simply does not refer to any global addresses, all variables are on the stack).

The `main.c` file provides definitions for the structs that match the ELF header and entries of the Program Header Table. So you can simply read the header out of the ELF file with the `open`, `lseek`, and `read` functions, read the offset of the program header table, get the number of the entries in the program header table from the ELF header, and read all entries one by one. If the entry has `ELF_PROG_LOAD` type, you will load it in memory.

To load the segment in memory, you can allocate executable memory with the following function

```
code_va = mmap(NULL, ph.memsz, PROT_READ | PROT_WRITE | PROT_EXEC, MAP_ANONYMOUS | MAP_PRIVATE, 0, 0);
```

where `ph.memsz` is the size of the segment that we currently are loading.

You should figure out where the entry point of your program is (it is in one of the segments, and since you map the segments at the addresses returned by the `mmap()` function, you need to find where it ends up in your address space.

If you found the entry point then type cast it to the function pointer that matches the function signature of the `sum` function and call it.

```
if(entry != NULL) {
    sum = entry;
    ret = sum(1, 2);
    printf("sum:%d\n", ret);
};
```

Your program should take the name of the executable elf file as its only argument and return the result produced by the function `main` present in that elf file. In particular, if the source code of the elf is:

```
unsigned int main(int a, int b) {
    return a + b;
}
```

Then your loader should print the following result:

```
$ ./main elf
sum:3
```

Please note that different elf files (one at the time) with different definitions of the function `main` will be passed to your loader in order to test the correctness of it.

Extra credit: (10% bonus)

Top

Try loading `elf-data.c` file. Can you explain why it crashes?

Fix the crash by either ensuring that the file is linked to work at the address that is available inside your process (it's a bit tricky to ensure that any specific address is available, but we'll accept any meaningful solution), or by performing a simple relocation step for the ELF file when it is loaded.

These [ELF tutorial](#) and [Executable and Linkable Format 101 Part 3: Relocations](#) resources can be helpful.

Submit your work

Top

Submit your solution (the zip archive) through Gradescope in the assignment called HW3 - The ELF Format . Please zip all of your files (main.c, Makefile) and submit them. If you have done extra credit then place that main.c and Makefile in an additional folder called "extra". **Please note that the non-optional part must be in the root of the zip archive, not inside yet another folder.**

You can resubmit as many times as you wish. If you have any problems with the structure the autograder will tell you. The structure of the zip file should be the following:

```
homework3.zip
├── main.c
├── Makefile
├── extra/ # optional
│   ├── main.c # optional
│   └── Makefile # optional
```